

# GMD 2A: MapReduce

Sabeur Aridhi

TELECOM Nancy, University of Lorraine

# Contents

- 1 Introduction
  - Big Data
  - Distributed computing
- 2 MapReduce
  - Programming model
  - Framework
  - Algorithms
- 3 Q and A

# Growth of data size

## Data sets

We need to analyze bigger and bigger datasets:

- **Web datasets:** web topology (Google: est. 46 billion pages)
- **Network datasets:** routing efficiency analysis, tcp-ip logs
- **Biological datasets:** genome sequencing data, protein patterns

## Big Data

Three main characteristics:

- **Volume:** Huge size of datasets (i.e., distributed storage).
- **Variety:** Several types of data such as text, image and video.
- **Velocity:** Data is generated at high speed (i.e., CERN experiments: one petabyte /sec).

## Key sources of Big Data

## Internet of Things (IoT)

- Networking of hardware equipment (e.g., household appliances, sensors, mobile phones)
- A significant source of Big Data



# Key sources of Big Data

## Smart cities

- Use of information technology to enhance quality, performance and interactivity of urban services in a city.
- Data is collected from sensors installed on:
  - ▶ utility poles
  - ▶ water lines
  - ▶ buses
  - ▶ trains
  - ▶ traffic lights, ...

# Key sources of Big Data - Smart cities



# Big Data in everyday life



# Big Data in everyday life



Does one minute fit into RAM? Five Minutes?



# Distributed Computing (DC)

- Requires a **Programming Model** (PM) that is inherently parallelizable.
- Requires a **framework** that runs the program and provides fault-tolerance and scalability.

# Distributed Computing (DC)

A **programming model** for DC should:

- abstract all low level details such as networking, scheduling, ...
- allow for wide range of functionality
- be easy to use

A **framework** for DC should provide:

- fault tolerance
- scalability
- data integrity

# Distributed File System

- Big Data requires storage over many machines
- Each computing node needs to access the data.
- A distributed file system (DFS) abstracts the distributed storage
- A widely used file system (FS) is Google DFS

# Distributed File System

- Big Data requires storage over many machines
- Each computing node needs to access the data.
- A distributed file system (DFS) abstracts the distributed storage
- A widely used file system (FS) is Google DFS

## Google's DFS

- data is replicated across the network. (Failure resistant, faster read operations).
- “namenode” stores file metadata e.g. name, location of chunks

# Distributed File System

- Big Data requires storage over many machines
- Each computing node needs to access the data.
- A distributed file system (DFS) abstracts the distributed storage
- A widely used file system (FS) is Google DFS

## Google's DFS

- data is replicated across the network. (Failure resistant, faster read operations).
- “namenode” stores file metadata e.g. name, location of chunks
- “datanode” stores chunks of that file

# Contents

- 1 Introduction
  - Big Data
  - Distributed computing
- 2 MapReduce
  - Programming model
  - Framework
  - Algorithms
- 3 Q and A

# Definition

Google, 2004:

*MapReduce refers to a programming model and the corresponding implementation for processing and generating large data sets.*

Created to help Google developers to analyze huge datasets

# Map and Reduce in MapReduce

- Data format: **key-value** pairs
- The user (you) must define two functions:

## Mapper

Map function:  $(k_1, v_1) \rightarrow \text{list}(k_2, v_2)$


## Reducer

Reduce function:  $(k_2, \text{list}(v_2)) \rightarrow \text{list}(v_3)$

- The Map functions are executed in parallel.
- $k_2, v_2$  are *intermediate* pairs.
- Mapper results are grouped using  $k_2$ .
- The Reduce functions are executed in parallel.



# MapReduce data flow (simplified)



## Basic steps of a MapReduce program:

- *Input is read from file system (fs)*
- The Map function is executed in parallel.
- *These results are written to fs*
- *Input is read from fs*
- The Reduce function is executed in parallel.
- *These results are written to fs*

## Example: Word Count (MapReduce's Hello World)

### Input

A set of documents stored in a DFS

### Output

The number of occurrences of each word in the set of documents.

## Example: Word Count (MapReduce's Hello World)

### Input

A set of documents stored in a DFS

### Output

The number of occurrences of each word in the set of documents.

### Idea

Work on each document in parallel and then “reduce” the results for each word.

## Example: Word Count (MapReduce's Hello World)

---

**Algorithm:** Map function

---

**Input:** String filename, String content

**foreach** word  $w$  **in** content **do**

$\text{EMITINTERMEDIATE}(w, 1);$

---

---

**Algorithm:** Reduce function

---

**Input:** String key, Iterator values

result  $\leftarrow 0;$

**foreach**  $v$  **in** values **do**

    result  $\leftarrow \text{result} + v;$

EMIT (key,result);

---

## Possible optimization

Creating a pair for each occurrence of a particular word is inefficient

---

**Algorithm:** Optimized Map function

---

**Input:** String filename, String content

$H \leftarrow \text{new HashTable};$

% H is an associative array that maps keys to values

**foreach** word w **in** content **do**

$H[w] \leftarrow H[w] + 1$

**foreach** word w **in** H **do**

$\text{EMITINTERMEDIATE}(w, H[w])$

---

# Combiner and Partitioner

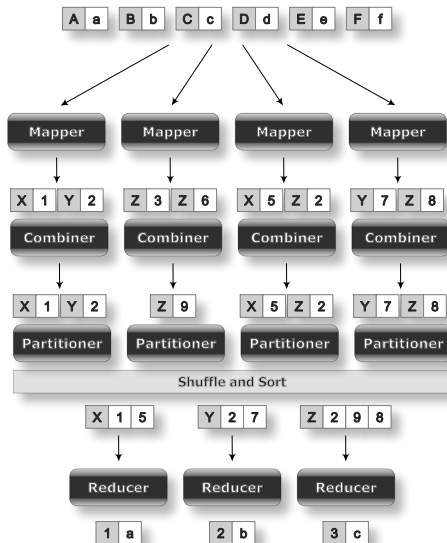
## Combiner

- Operates between Mapper and Reducer.
- *Combiner* :  $(k_2, \text{list}(v_2)) \rightarrow (k_2, \text{list}(v_3))$
- The combiner is applied before the global sort.
- Just like in the WordCount example.

## Partitioner

Before global sort, decide which reducer receives which key.

# Complete MapReduce flow






# Fault tolerance management

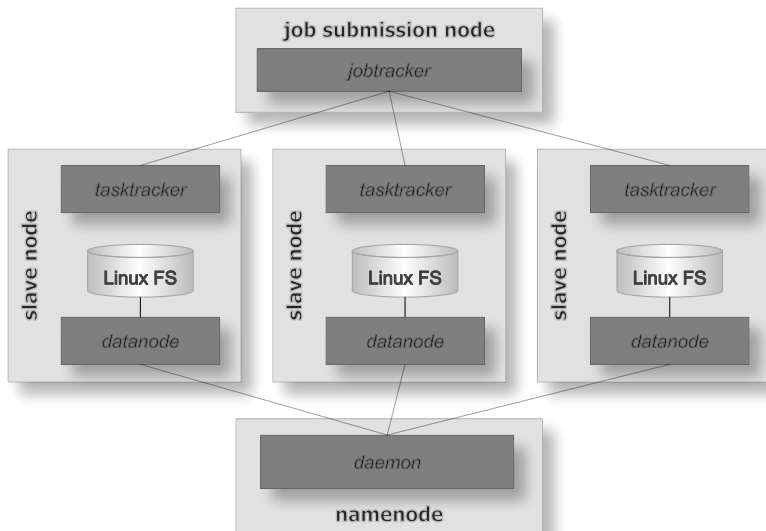
- In case of a worker failure
  - ▶ map (and not reduce) tasks completed reset to idle state
  - ▶ map and reduce tasks in progress reset to idle
  - ▶ Notification to and re-execution by workers executing reduce tasks
- in case of Master failure
  - ▶ Version of 2004:
    - ★ Abort the computation and retry
  - ▶ In MRv2 (YARN):
    - ★ Periodic check points of the data structures
    - ★ Launching of a new copy from the check point

# Framework

- Google's MapReduce framework is not freely available
- Open source alternative: Apache's Hadoop
- Offered as service in Amazon Elastic Computing (via virtual machines)



# Hadoop Architecture



# Hadoop Details

- Library in Java
- Started by Yahoo in 2006
- By default, Map and Reduce functions must be written in Java
- Possibility to use external programs as mapper and reducers

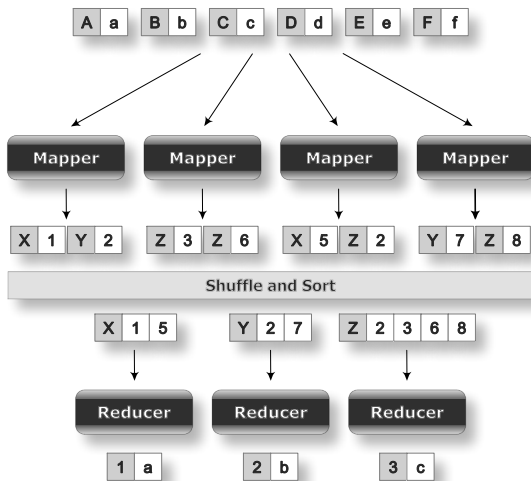
# Inverted Index (Job interview question ;)

## Problem

We are given a set of webpages  $P = \{P_1, \dots, P_n\}$ .

We want to know for each word in the dataset, in which documents it appears (and how many times)

# Reminder of MapReduce flow



## Baseline solution

---

**Algorithm:** Map function(String filename, String content)

---

```
 $H \leftarrow \text{new HashTable};$   
foreach word  $w$  in content do  
   $H[w] \leftarrow H[w] + 1$   
foreach word  $w$  in  $H$  do  
   $\text{EMITINTERMEDIATE}(w, (\text{filename}, H[w]))$ 
```

---

---

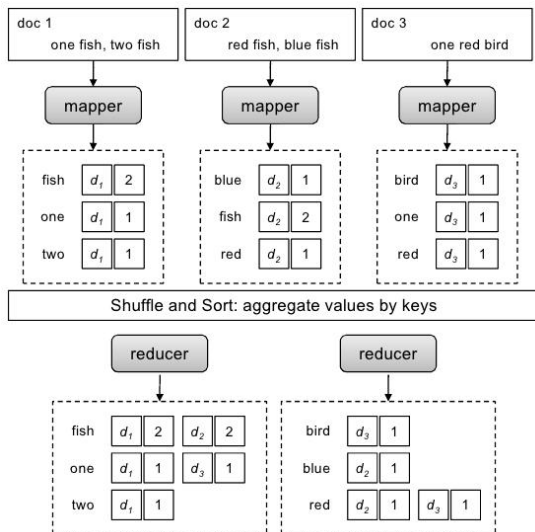
**Algorithm:** Reduce function(String key, Iterator values)

---

```
result  $\leftarrow []$ ;  
foreach  $v$  in values do  
  result.add( $v$ );  
result.sortbyfreq();  
 $\text{EMIT}(\text{key}, \text{result});$ 
```

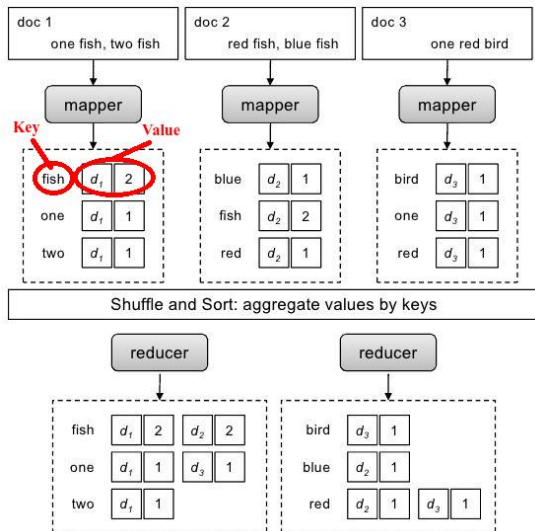
---

# Simple run





# Simple run



# Contents

- 1 Introduction
  - Big Data
  - Distributed computing
- 2 MapReduce
  - Programming model
  - Framework
  - Algorithms
- 3 Q and A

That's it. Why was this important?



Questions?  
sabeur.aridhi@loria.fr